

## 1. Introduction to NumPy

### What is NumPy?

- **NumPy (Numerical Python)** is a powerful Python library used for numerical and scientific computing.
- It provides a **multidimensional array object** called `ndarray`, which is much faster and more efficient than Python's built-in lists.
- NumPy is the foundation for many other libraries in **data science, machine learning, and scientific computing** (such as Pandas, SciPy, Scikit-learn, TensorFlow, PyTorch).

👉 In short: NumPy = **fast math with arrays + a huge toolbox for linear algebra, statistics, and random numbers.**

### Why NumPy vs Python Lists?

#### 1. Speed

- NumPy is written in **C** under the hood, so operations are executed much faster than pure Python lists.

Example:

```
import numpy as np
import time

# Using Python list
python_list = list(range(1_000_000))
start = time.time()
[ x*2 for x in python_list ]
print("Python list time:", time.time() - start)

# Using NumPy array
numpy_array = np.arange(1_000_000)
start = time.time()
numpy_array * 2
print("NumPy array time:", time.time() - start)
```

👉 NumPy will be **10–100x faster** depending on the operation.

#### 2. Memory Efficiency

- A Python list stores each element as a **full Python object** with extra overhead.
- A NumPy array stores data in a **continuous block of memory** with fixed data types, making it much more memory efficient.

Example:

```
import numpy as np
import sys

# Python list of 1 million integers
python_list = list(range(1_000_000))
print("Size of Python list:", sys.getsizeof(python_list))

# NumPy array of 1 million integers
numpy_array = np.arange(1_000_000)
print("Size of NumPy array:", numpy_array.nbytes)
```

👉 NumPy arrays take **much less memory**.

#### 3. Functionality

- Python lists only support basic operations like `append`, `insert`, etc.
- NumPy arrays support:
  - **Vectorized operations** (e.g., `array * 2` works on all elements at once)

- **Mathematical functions** (`sin`, `log`, `exp`, etc.)
- **Linear algebra** (`dot`, `inv`, eigenvalues)
- **Random numbers**
- **Statistics** (mean, median, variance, etc.)

Example:

```
import numpy as np

arr = np.array([1, 2, 3, 4, 5])

print("Multiply by 2:", arr * 2)
print("Square:", arr ** 2)
print("Mean:", np.mean(arr))
print("Sine:", np.sin(arr))
```

## ◆ Installing NumPy

You can install NumPy in your Python environment using **pip**:

```
pip install numpy
```

If you are using **Anaconda**, NumPy usually comes pre-installed. If not, install via:

```
conda install numpy
```

Then verify installation in Python:

```
import numpy as np
print(np.__version__)
```

Start coding or [generate](#) with AI.

## ✓ ■ 2. NumPy Basics (Full Guide)

### ◆ 2.1 Importing NumPy

```
import numpy as np
```

⚡ **Tip:** Always use `np` as alias (standard in the Python community).

### ◆ 2.2 Creating Arrays

#### 1. From Lists and Tuples

```
arr = np.array([1, 2, 3, 4, 5]) # 1D array
matrix = np.array([[1, 2], [3, 4]]) # 2D array
```


- NumPy will **infer dtype automatically**.
- Override with `dtype`:

```
arr = np.array([1, 2, 3], dtype=np.float32)
```

⚡ **Tip:** Use `dtype` wisely for memory efficiency in big data.

#### 2. Predefined Arrays

```
zeros = np.zeros((2, 3))    # All zeros
ones = np.ones((3, 2))     # All ones
full = np.full((2, 2), 9)   # All filled with 9
```

 **Tip:** Faster for initializing matrices before computations.

### 3. Identity and Diagonal

```
I = np.eye(3)              # Identity matrix
diag = np.diag([5, 10, 15]) # Custom diagonal
```

 **Tip:** Very useful for **linear algebra**.

### 4. Range-Based Arrays


```
a = np.arange(0, 10, 2)    # [0,2,4,6,8]
b = np.linspace(0, 1, 5)   # 5 values evenly spaced
c = np.logspace(1, 3, 3)   # [10, 100, 1000]
```

 **Tip:** Use `linspace` in plotting & simulation (smooth values).

### 5. Random Arrays


```
np.random.seed(42) # reproducibility

uniform = np.random.rand(2, 3)    # Uniform [0,1)
normal = np.random.randn(2, 3)    # Standard normal
integers = np.random.randint(1, 10, size=(3, 3)) # 1-9
choice = np.random.choice([10, 20, 30], size=5) # Random pick
```

 **Tip:** Always fix `seed` in experiments for **reproducible results**.

### 6. Special Arrays


```
empty = np.empty((2, 2)) # Random memory garbage
identity = np.identity(4)
```

 **Warning:** `np.empty()` is fast but contains **uninitialized values**.

## ◆ 2.3 Array Attributes

```
arr = np.arange(12).reshape(3, 4)

print("Array:\n", arr)
print("Shape:", arr.shape)    # (3,4)
print("Dimensions:", arr.ndim) # 2
print("Size:", arr.size)      # 12
print("Data type:", arr.dtype)
print("Item size:", arr.itemsize, "bytes")
print("Total memory:", arr.nbytes, "bytes")
```

 **Tip:** Use `arr.shape` unpacking:

```
rows, cols = arr.shape
```

## ◆ 2.4 Reshaping Arrays

```
arr = np.arange(12)          # 0..11
reshaped = arr.reshape(3, 4)

print("Reshaped:\n", reshaped)
```

Flattening:

```
print("Flatten:", reshaped.flatten()) # copy
print("Ravel:", reshaped.ravel())    # view
```

⚡ **Tip:** Use `reshape(-1)` to flatten quickly.

## ◆ 2.5 Copy vs View

```
a = np.array([1, 2, 3])
b = a.view()   # view (linked to original)
c = a.copy()   # copy (independent)

a[0] = 99
print("Original:", a) # [99, 2, 3]
print("View:", b)     # [99, 2, 3]
print("Copy:", c)     # [1, 2, 3]
```

⚡ **Tip:** Always use `.copy()` if you don't want accidental modifications.

## ◆ 2.6 Data Type Casting

```
arr = np.array([1.5, 2.8, 3.9])
int_arr = arr.astype(int)
print(int_arr) # [1, 2, 3]
```

⚡ **Tip:** `.astype()` is critical in ML preprocessing (float32 for GPUs).

## ■ Examples (Zero → Hero)

### ✅ Beginner: Create & Inspect

```
arr = np.arange(1, 10).reshape(3, 3)
print("Array:\n", arr)
print("Mean:", np.mean(arr))
print("Sum of all elements:", arr.sum())
```

### ✅ Intermediate: Dice Simulator

```
np.random.seed(0)
dice = np.random.randint(1, 7, size=20)
print("Dice rolls:", dice)
print("Frequency of 6:", np.count_nonzero(dice == 6))
```

### ✅ Advanced: Monte Carlo $\pi$

```
np.random.seed(42)
N = 1_000_000
x, y = np.random.rand(N), np.random.rand(N)

inside = (x**2 + y**2) <= 1
pi_est = 4 * np.sum(inside) / N
```

```
print("Estimated  $\pi$ :", pi_est)
```

## ✓ Hero: Gradient Image

```
import matplotlib.pyplot as plt

gradient = np.linspace(0, 1, 256)
image = np.outer(gradient, gradient)

plt.imshow(image, cmap="viridis")
plt.colorbar()
plt.title("Generated Gradient")
plt.show()
```

## Exercises

### Beginner

1. Create a 1D array from 10 to 50 with step 5.
2. Create a 3x3 array filled with True (boolean dtype).
3. Create an array of 100 random numbers between 0 and 1 and find its mean.

### Intermediate

4. Create a 4x4 identity matrix and replace the diagonal with numbers 1–4.
5. Create a 5x5 array with values from 0 to 24, then reshape it into 25x1.
6. Generate 10 random integers between 1 and 100. Find the max, min, and mean.

### Advanced

7. Create a 6x6 checkerboard pattern (0s and 1s alternating).
8. Generate 1 million random numbers (normal distribution) and calculate mean & std.
9. Simulate tossing 2 dice 1000 times and estimate the probability of getting a sum of 7.

Start coding or [generate](#) with AI.

## ✓ Exercises with Solutions

### Beginner

#### 1. Create a 1D array from 10 to 50 with step 5.

```
import numpy as np

arr = np.arange(10, 55, 5)
print(arr)  # [10 15 20 25 30 35 40 45 50]
```

✓ **Tip:** `np.arange(start, stop, step)` is like Python `range()` but returns an array.

#### 2. Create a 3x3 array filled with True (boolean dtype).

```
arr = np.ones((3, 3), dtype=bool)
print(arr)
```

✓ **Tip:** `np.ones + dtype=bool` is a quick trick for boolean masks.

#### 3. Create an array of 100 random numbers between 0 and 1 and find its mean.

```
arr = np.random.rand(100)
print("Mean:", arr.mean())
```

✓ **Tip:** NumPy arrays have built-in aggregation methods ( `.mean()` , `.sum()` ).

## Intermediate

### 4. Create a 4x4 identity matrix and replace the diagonal with numbers 1–4.

```
arr = np.eye(4)
arr[np.arange(4), np.arange(4)] = [1, 2, 3, 4]
print(arr)
```

✓ **Tip:** Use `arr[np.arange(n), np.arange(n)]` to select diagonals efficiently.

### 5. Create a 5x5 array with values from 0 to 24, then reshape it into 25x1.

```
arr = np.arange(25).reshape(5, 5)
print("5x5:\n", arr)

reshaped = arr.reshape(25, 1)
print("25x1:\n", reshaped)
```

✓ **Tip:** `reshape(-1, 1)` is a shortcut to flatten into a column vector.

### 6. Generate 10 random integers between 1 and 100. Find the max, min, and mean.

```
arr = np.random.randint(1, 101, 10)
print("Numbers:", arr)
print("Max:", arr.max())
print("Min:", arr.min())
print("Mean:", arr.mean())
```

✓ **Tip:** NumPy has `arr.max()`, `arr.min()`, `arr.mean()` optimized in C.

## Advanced

### 7. Create a 6x6 checkerboard pattern (0s and 1s alternating).

```
arr = np.zeros((6, 6), dtype=int)
arr[1::2, ::2] = 1 # odd rows, even cols
arr[:, 1::2] = 1 # even rows, odd cols
print(arr)
```

✓ **Tip:** Use **slicing tricks** ( `::2` ) for pattern generation.

### 8. Generate 1 million random numbers (normal distribution) and calculate mean & std.

```
arr = np.random.randn(1_000_000)
print("Mean:", arr.mean())
print("Standard Deviation:", arr.std())
```

✓ **Tip:** For large data, NumPy is **much faster** than Python's `statistics` module.

### 9. Simulate tossing 2 dice 1000 times and estimate the probability of getting a sum of 7.

```
dice1 = np.random.randint(1, 7, 1000)
dice2 = np.random.randint(1, 7, 1000)
sums = dice1 + dice2
```

```
prob_7 = np.count_nonzero(sums == 7) / 1000
print("Estimated Probability of 7:", prob_7)
```

✔ **Tip:** Vectorized operations (`dice1 + dice2`) make this super efficient.

## Summary (NumPy Basics Module)

By now you've learned: ✔ How to create arrays (zeros, ones, random, ranges, identity, etc.) ✔ How to check attributes (shape, size, dtype, memory) ✔ How to reshape, flatten, copy, and cast arrays ✔ Real applications (dice rolls, Monte Carlo, image generation) ✔ Solved exercises from beginner to advanced

Start coding or [generate](#) with AI.

## ▼ Section 3: Array Attributes in NumPy

When working with NumPy, understanding the **properties of arrays** is essential. These attributes help you **inspect, reshape, and manipulate arrays efficiently**.

### ◆ 3.1 Basic Attributes

Let's create a sample array:

```
import numpy as np

arr = np.array([[1, 2, 3], [4, 5, 6]])
print(arr)
```

Output:

```
[[1 2 3]
 [4 5 6]]
```

#### 1. Shape ( `.shape` )

Returns the **dimensions** of the array as a tuple.

```
print("Shape:", arr.shape) # (2, 3)
```

✔ This means 2 rows × 3 columns.

#### 2. Size ( `.size` )

Returns the **total number of elements**.

```
print("Size:", arr.size) # 6
```

#### 3. Dimensions ( `.ndim` )

Returns the number of **dimensions (axes)**.

```
print("Dimensions:", arr.ndim) # 2
```

#### 4. Data Type ( `.dtype` )

Returns the **type of elements** in the array.

```
print("Data Type:", arr.dtype) # int64 (depends on system)
```

## 5. Item Size & Total Memory

```
print("Item size (bytes):", arr.itemsize)
print("Total size (bytes):", arr.nbytes)
```

✅ Handy for memory optimization.

⚡ **Tip:** You can convert data type with `arr.astype(new_dtype)`

```
float_arr = arr.astype(float)
print(float_arr.dtype) # float64
```

## ◆ 3.2 Reshaping Arrays

Sometimes, you want to change the **shape** without changing data.

### 1. reshape()

```
arr = np.arange(12)
reshaped = arr.reshape(3, 4)
print(reshaped)
```

Output:

```
[[ 0  1  2  3]
 [ 4  5  6  7]
 [ 8  9 10 11]]
```

⚡ **Trick:** Use `-1` to let NumPy infer the dimension.

```
print(arr.reshape(2, -1)) # 2 rows, NumPy auto-calculates columns
```

### 2. ravel()

Flattens array into **1D view** (does not copy data).

```
flat_view = reshaped.ravel()
flat_view[0] = 99
print("Flat view:", flat_view)
print("Original:", reshaped) # Also changes!
```

### 3. flatten()

Flattens array into **1D copy** (independent).

```
flat_copy = reshaped.flatten()
flat_copy[0] = -1
print("Flat copy:", flat_copy)
print("Original:", reshaped) # Not affected
```

⚡ **Rule of Thumb:**

- `ravel()` → faster, returns a **view** (changes affect original).
- `flatten()` → safer, returns a **copy** (original stays intact).

## ◆ 3.3 Practical Examples

### Example 1: Image Reshaping



If you have a grayscale image with 28×28 pixels:

```
img = np.arange(784).reshape(28, 28)
print(img.shape) # (28, 28)

flat_img = img.ravel()
print(flat_img.shape) # (784,)
```

✓ Used in **deep learning (MNIST dataset)** where each image must be flattened.

## Example 2: Converting Sensor Data

Suppose you receive a **stream of 1D sensor values** (100 readings):

```
data = np.arange(100)
reshaped = data.reshape(10, 10) # group into 10 samples × 10 features
```

✓ This way, machine learning models can process it in batches.

## ◆ 3.4 Exercises

### Beginner

1. Create a 1D array of 15 elements and check: shape, size, ndim, and dtype.
2. Convert an array of integers into float type.

### Intermediate

3. Create a 3×4 array and reshape it into:
  - (2, 6)
  - (6, 2, 1)
4. Flatten an array using both `ravel()` and `flatten()`. Show the difference by modifying one element.

### Advanced

5. Create a 3D array of shape (2, 3, 4).
  - Find its size, ndim, and reshape it into (4, 6).
6. Generate a (100, ) array of random values. Reshape it into (10, 10) and calculate the mean of each row.

## ◆ 3.5 Solutions

### Beginner

```
arr = np.arange(15)
print(arr.shape) # (15,)
print(arr.size) # 15
print(arr.ndim) # 1
print(arr.dtype) # int64
```

```
float_arr = arr.astype(float)
print(float_arr.dtype) # float64
```

### Intermediate

```
arr = np.arange(12).reshape(3, 4)
print(arr.reshape(2, 6))
print(arr.reshape(6, 2, 1))
```

```
flat_view = arr.ravel()
flat_copy = arr.flatten()

flat_view[0] = 999
print("Original after ravel change:\n", arr)

flat_copy[0] = -1
print("Original after flatten change:\n", arr)
```

## Advanced

```
arr = np.arange(24).reshape(2, 3, 4)
print("Size:", arr.size) # 24
print("Dimensions:", arr.ndim) # 3
```

```
reshaped = arr.reshape(4, 6)
print(reshaped.shape)
```

```
data = np.random.rand(100).reshape(10, 10)
row_means = data.mean(axis=1)
print("Row means:", row_means)
```

✅ With this, you now master:

- **Array inspection (shape, size, ndim, dtype)**
- **Reshaping arrays safely**
- **Flattening (ravel vs flatten)**
- **Practical use-cases (ML, images, sensors)**

Start coding or [generate](#) with AI.

## 3. Array Attributes in NumPy

NumPy arrays (`ndarray`) come with built-in **attributes** that describe their structure. These attributes tell us about **dimensions, shape, size, and data type**.

### 3.1 Shape, Size, ndim, and dtype

#### ✅ shape

- Returns a **tuple** showing the number of elements in each dimension.

```
import numpy as np

arr = np.array([[1, 2, 3], [4, 5, 6]])
print("Array:\n", arr)
print("Shape:", arr.shape) # (2 rows, 3 columns)
```

#### ✅ size

- Returns the **total number of elements** in the array.

```
print("Size:", arr.size) # 6 elements in total
```

#### ✅ ndim

- Returns the **number of dimensions** (axes) of the array.

```
print("Dimensions:", arr.ndim) # 2 (rows and columns)
```

## ✔ dtype

- Shows the **data type** of array elements.

```
print("Data type:", arr.dtype) # int64 (on most systems)
```

👉 You can also **force a dtype** when creating arrays:

```
arr_float = np.array([1, 2, 3], dtype=np.float32)
print("Forced dtype:", arr_float.dtype)
```

## ◆ 3.2 Reshaping Arrays

Reshaping allows you to **change the structure** of an array without changing its data.

## ✔ reshape()

- Creates a new array with a new shape (must have same total elements).

```
arr = np.arange(12) # [0,1,2,...,11]
reshaped = arr.reshape(3, 4)
print("Reshaped (3x4):\n", reshaped)
```

👉 Tip: Use `-1` as a placeholder to let NumPy calculate automatically.

```
arr = np.arange(12)
reshaped = arr.reshape(2, -1) # 2 rows, auto columns
print("Reshaped with -1:\n", reshaped)
```

## ✔ ravel()

- Returns a **flattened 1D array**, but gives a **view (reference)**, not a copy.
- Changing the flattened array may affect the original.

```
arr = np.array([[1, 2], [3, 4]])
ravelled = arr.ravel()
print("Ravelled:", ravelled)
ravelled[0] = 99
print("Modified original:\n", arr) # arr changes too
```

## ✔ flatten()

- Returns a **copy** of the array in 1D form.
- Changing it does **not** affect the original.

```
arr = np.array([[1, 2], [3, 4]])
flattened = arr.flatten()
flattened[0] = 99
print("Flattened:", flattened)
print("Original stays same:\n", arr)
```

## ◆ 3.3 Tips & Tricks

- ✔ Use `.shape` before reshaping to avoid mismatches.
- ✔ Prefer `ravel()` over `flatten()` for efficiency (when you don't need a copy).
- ✔ Use `reshape(-1)` to quickly flatten arrays.
- ✔ Use `.ndim` when writing general functions to handle arrays of any dimension.

### ◆ 3.4 Examples from Zero to Hero

#### Example 1: Reshape for ML Dataset

Suppose you load 28×28 images (like MNIST dataset) and need to flatten them into a 1D vector for training:

```
images = np.arange(28*28*10).reshape(10, 28, 28) # 10 images
print("Original shape:", images.shape) # (10, 28, 28)

flattened = images.reshape(10, -1) # (10, 784)
print("Flattened shape:", flattened.shape)
```

#### Example 2: Converting 1D → 2D Matrix for Math

```
arr = np.arange(9)
matrix = arr.reshape(3, 3)
print("Matrix:\n", matrix)
```

#### Example 3: Understanding Copy vs View

```
arr = np.array([1, 2, 3, 4])

# Ravel (view)
r = arr.ravel()
r[0] = 100
print("After ravel modification:", arr)

# Flatten (copy)
f = arr.flatten()
f[1] = 200
print("After flatten modification:", arr) # Original unchanged
```

### ◆ 3.5 Exercises

💡 Try these exercises to master attributes:

1. Create a 5x5 NumPy array of numbers 0–24. Print its `shape`, `size`, `ndim`, and `dtype`.
2. Reshape it into a 25x1 column vector.
3. Use `ravel()` to flatten it, then modify one element and check if the original changed.
4. Do the same with `flatten()` and compare results.
5. Create a random array of shape (2, 3, 4). Reshape it into (4, 6) using `-1`.

✅ Now you know **array attributes and reshaping tricks** — essential for any data science or ML workflow.

Start coding or [generate](#) with AI.

## ✓ ■ 4. Indexing and Slicing in NumPy

NumPy arrays are **indexed and sliced** similar to Python lists, but with much more power.

### ◆ 1D Array Indexing

```
import numpy as np

arr = np.array([10, 20, 30, 40, 50])
print(arr[0]) # First element -> 10
print(arr[-1]) # Last element -> 50
print(arr[2]) # Third element -> 30
```

- ✔ Works the same as Python lists.

## ◆ 2D Array Indexing

- For **2D arrays**, use [row, column] format.

```
arr2d = np.array([[10, 20, 30],
                  [40, 50, 60],
                  [70, 80, 90]])

print(arr2d[0, 0]) # 10 (1st row, 1st column)
print(arr2d[1, 2]) # 60 (2nd row, 3rd column)
print(arr2d[-1, -1]) # 90 (last row, last column)
```

## ◆ 3D Array Indexing

- For **3D arrays**, indexing goes [depth, row, column].

```
arr3d = np.array([[[1, 2], [3, 4]],
                  [[5, 6], [7, 8]]])

print(arr3d[0, 0, 1]) # 2 (1st block, 1st row, 2nd col)
print(arr3d[1, 1, 0]) # 7 (2nd block, 2nd row, 1st col)
```

## ◆ Slicing Arrays (start:stop:step)

1D Example:

```
arr = np.array([0, 10, 20, 30, 40, 50, 60])

print(arr[1:5])      # [10 20 30 40] (from index 1 to 4)
print(arr[:4])       # [0 10 20 30] (from start to 3)
print(arr[::2])      # [0 20 40 60] (every 2nd element)
print(arr[::-1])     # [60 50 40 30 20 10 0] (reversed)
```

2D Example:

```
arr2d = np.array([[10, 20, 30, 40],
                  [50, 60, 70, 80],
                  [90, 100, 110, 120]])

print(arr2d[0:2, 1:3]) # [[20 30], [60 70]]
print(arr2d[:, 2])     # [30 70 110] (entire 3rd column)
print(arr2d[1, :])     # [50 60 70 80] (entire 2nd row)
```

## ◆ Negative Indexing

- NumPy supports negative indexing to count from the end.

```
arr = np.array([10, 20, 30, 40, 50])
print(arr[-1])      # 50
print(arr[-3:])     # [30 40 50]
print(arr[:-2])     # [10 20 30]
```

## ◆ Fancy Indexing (using lists/arrays)

- Fancy indexing lets you extract **multiple elements at once** using a list/array of indices.

```
arr = np.array([10, 20, 30, 40, 50])

print(arr[[0, 2, 4]]) # [10 30 50]
print(arr[[1, -1]])   # [20 50]
```

## 2D Example:

```
arr2d = np.array([[10, 20, 30],
                  [40, 50, 60],
                  [70, 80, 90]])

print(arr2d[[0, 2], [1, 2]])
# Picks (0,1) → 20 and (2,2) → 90 → [20 90]
```

## ◆ Boolean Indexing

- Boolean indexing is super powerful for filtering arrays.

```
arr = np.array([10, 20, 30, 40, 50])

print(arr[arr > 25]) # [30 40 50]
print(arr[arr % 20 == 0]) # [20 40]
```

## With 2D Arrays:

```
arr2d = np.array([[5, 10, 15],
                  [20, 25, 30],
                  [35, 40, 45]])

print(arr2d[arr2d > 20]) # [25 30 35 40 45]
```

## ⚡ Tips & Tricks

1. Use **slicing with steps** (`arr[: :2]`) instead of loops.
2. Boolean indexing is great for **filtering** datasets (`arr[arr > threshold]`).
3. Fancy indexing can select **non-contiguous elements** efficiently.
4. `np.where(condition)` can combine Boolean and fancy indexing.

### Example:

```
arr = np.array([5, 10, 15, 20, 25])
indices = np.where(arr > 12)
print(indices) # (array([2, 3, 4]),)
print(arr[indices]) # [15 20 25]
```

## 📝 Exercises

### Beginner

1. Create an array `[1, 2, 3, 4, 5, 6, 7, 8, 9]` and extract:
  - The last 3 elements
  - Every second element
  - The array reversed
2. Create a 3×3 array of numbers 1–9. Extract:
  - The 2nd row
  - The 1st column
  - The bottom-right element

## Intermediate

3. Create a 4×4 array of numbers 0–15. Extract:

- The submatrix `[[5, 6], [9, 10]]`
- All even numbers

4. Using Boolean indexing, replace all numbers greater than 10 with -1.

## Advanced (Hero 🦸)

5. Create a 5×5 array with numbers from 0–24.

- Extract the diagonal using fancy indexing.
- Extract all border elements (top row, bottom row, first col, last col).

6. Given an array `[3, 6, 9, 12, 15, 18, 21]`:

- Use `np.where` to find indices of numbers divisible by 9.
- Replace those numbers with 99.

Start coding or [generate](#) with AI.

## 4. Indexing and Slicing Exercises with Solutions

### ◆ 1. 1D Array Indexing

**Q1.** Create a 1D NumPy array of numbers 10 to 19. Print:

- the 1st element
- the last element
- elements from index 2 to 5

```
import numpy as np

arr = np.arange(10, 20) # [10 11 12 13 14 15 16 17 18 19]

print("Array:", arr)

# First element
print("First element:", arr[0])

# Last element
print("Last element:", arr[-1])

# Elements from index 2 to 5
print("Elements 2 to 5:", arr[2:6])
```

#### ✅ Output:

```
Array: [10 11 12 13 14 15 16 17 18 19]
First element: 10
Last element: 19
Elements 2 to 5: [12 13 14 15]
```

### ◆ 2. 2D Array Indexing

**Q2.** Create a 3×4 array with values from 1 to 12. Print:

- element at row 1, col 2
- entire 2nd row
- entire 3rd column

```
arr2d = np.arange(1, 13).reshape(3, 4)

print("2D Array:\n", arr2d)

# Element at row 1, col 2
print("Element (row=1, col=2):", arr2d[1, 2])

# Entire 2nd row
print("Second row:", arr2d[1])

# Entire 3rd column
print("Third column:", arr2d[:, 2])
```

#### ✓ Output:

```
2D Array:
[[ 1  2  3  4]
 [ 5  6  7  8]
 [ 9 10 11 12]]

Element (row=1, col=2): 7
Second row: [5 6 7 8]
Third column: [ 3  7 11]
```

### ◆ 3. 3D Array Indexing

**Q3.** Create a 2×2×3 array with numbers 1–12. Print:

- first 2D block
- element at block 1, row 0, col 2

```
arr3d = np.arange(1, 13).reshape(2, 2, 3)

print("3D Array:\n", arr3d)

# First 2D block
print("First block:\n", arr3d[0])

# Element at block=1, row=0, col=2
print("Element [1,0,2]:", arr3d[1, 0, 2])
```

#### ✓ Output:

```
3D Array:
[[[ 1  2  3]
 [ 4  5  6]]

 [[ 7  8  9]
 [10 11 12]]]

First block:
[[1 2 3]
 [4 5 6]]
Element [1,0,2]: 9
```

### ◆ 4. Slicing

**Q4.** For array `[0,1,2,3,4,5,6,7,8,9]`:

- elements from index 2 to 7 (exclusive)
- every 2nd element
- reverse the array



```
arr = np.arange(10)

print("Array:", arr)
print("Index 2 to 7:", arr[2:8])
print("Every 2nd element:", arr[::2])
print("Reversed:", arr[::-1])
```

✔ **Output:**

```
Array: [0 1 2 3 4 5 6 7 8 9]
Index 2 to 7: [2 3 4 5 6 7]
Every 2nd element: [0 2 4 6 8]
Reversed: [9 8 7 6 5 4 3 2 1 0]
```

## ◆ 5. Negative Indexing

Q5. From [10, 20, 30, 40, 50] print:

- last 3 elements using negative slicing
- last element

```
arr = np.array([10, 20, 30, 40, 50])

print("Last 3 elements:", arr[-3:])
print("Last element:", arr[-1])
```

✔ **Output:**

```
Last 3 elements: [30 40 50]
Last element: 50
```

## ◆ 6. Fancy Indexing

Q6. From [10,20,30,40,50], select elements at positions 0, 2, and 4.

```
arr = np.array([10, 20, 30, 40, 50])
indices = [0, 2, 4]
print("Selected elements:", arr[indices])
```

✔ **Output:**

```
Selected elements: [10 30 50]
```

## ◆ 7. Boolean Indexing

Q7. From [1,2,3,4,5,6,7,8,9,10], select:

- even numbers
- numbers greater than 5

```
arr = np.arange(1, 11)

print("Even numbers:", arr[arr % 2 == 0])
print("Greater than 5:", arr[arr > 5])
```

✔ **Output:**

```
Even numbers: [ 2  4  6  8 10]
Greater than 5: [ 6  7  8  9 10]
```

## Practice Exercises (Try Yourself)

1. Create a 2D array `5x5` with numbers 1–25. Extract:
  - the center element
  - the first row
  - the last column
2. Create a 3D array `3x3x3` with numbers 1–27. Extract:
  - the last block
  - the diagonal elements of the first block
3. From `[5, 10, 15, 20, 25, 30, 35, 40]`, get:
  - numbers divisible by 10
  - reverse using slicing

Start [coding](#) or [generate](#) with AI.

## ✓ 5. Array Operations in NumPy

### 1. Element-wise Operations

NumPy supports **element-wise arithmetic** on arrays of the same shape.

```
import numpy as np

a = np.array([1, 2, 3, 4])
b = np.array([10, 20, 30, 40])

print(a + b)  # [11 22 33 44]
print(a - b)  # [-9 -18 -27 -36]
print(a * b)  # [10 40 90 160]
print(b / a)  # [10. 10. 10. 10.]
print(a ** 2) # [ 1  4  9 16]
```

✓ **Tip:** NumPy automatically **vectorizes** operations → much faster than Python loops.

### 2. Universal Functions (ufuncs)

NumPy provides **fast mathematical functions** (ufuncs) that work element-wise.

```
arr = np.array([1, 4, 9, 16, 25])

print(np.sqrt(arr))  # Square root [1. 2. 3. 4. 5.]
print(np.exp(arr))   # Exponential
print(np.log(arr))   # Natural log
print(np.log10(arr)) # Log base 10
print(np.sin(arr))   # Trigonometric functions
print(np.cos(arr))
```

✓ **Tip:** Ufuncs are highly optimized in C → thousands of times faster than manual calculations.

### 3. Aggregations

NumPy allows **aggregation functions** to summarize arrays.

```
data = np.array([1, 2, 3, 4, 5])

print(data.sum())  # 15
print(data.mean()) # 3.0
print(data.max())  # 5
```

```
print(data.min()) # 1
print(data.std()) # Standard deviation
print(data.var()) # Variance
```

## 4. Axis Operations

For multi-dimensional arrays, operations can be applied along **rows or columns**.

```
matrix = np.array([[1, 2, 3],
                   [4, 5, 6],
                   [7, 8, 9]])

print(matrix.sum(axis=0)) # Sum by columns [12 15 18]
print(matrix.sum(axis=1)) # Sum by rows   [6 15 24]

print(matrix.mean(axis=0)) # Column mean [4. 5. 6.]
print(matrix.mean(axis=1)) # Row mean [2. 5. 8.]
```

✓ **Tip:** `axis=0` → column-wise, `axis=1` → row-wise.

## Exercises

### Q1. Element-wise Operations

Given:

```
x = np.array([2, 4, 6, 8])
y = np.array([1, 3, 5, 7])
```

1. Add `x` and `y`.
2. Subtract `y` from `x`.
3. Multiply `x` and `y`.
4. Compute `x / y`.
5. Square all elements of `x`.

### Q2. Universal Functions

Given:

```
arr = np.array([0, np.pi/2, np.pi])
```

1. Compute `sin(arr)`.
2. Compute `cos(arr)`.
3. Compute `tan(arr)`.

### Q3. Aggregations

Given:

```
data = np.array([10, 20, 30, 40, 50])
```

1. Find the sum, mean, max, min, std, var.

### Q4. Axis Operations

Given:

```
mat = np.array([[2, 4, 6],
                [1, 3, 5],
                [7, 9, 11]])
```

1. Find column sums.
2. Find row means.
3. Find max in each row.

## ✓ Solutions

### Solution Q1

```
print(x + y) # [ 3  7 11 15]
print(x - y) # [1 1 1 1]
print(x * y) # [ 2 12 30 56]
print(x / y) # [2.  1.33333333 1.2  1.14285714]
print(x ** 2) # [ 4 16 36 64]
```

### Solution Q2

```
print(np.sin(arr)) # [0.  1.  0.]
print(np.cos(arr)) # [ 1.  0. -1.]
print(np.tan(arr)) # [ 0.  1.63312394e+16  0.] (= infinity at pi/2)
```

### Solution Q3

```
print(data.sum()) # 150
print(data.mean()) # 30.0
print(data.max()) # 50
print(data.min()) # 10
print(data.std()) # 14.1421
print(data.var()) # 200.0
```

### Solution Q4

```
print(mat.sum(axis=0)) # [10 16 22]
print(mat.mean(axis=1)) # [4.  3.  9.]
print(mat.max(axis=1)) # [ 6  5 11]
```

Start coding or [generate](#) with AI.

## ✓ 6. Broadcasting in NumPy

### What is Broadcasting?

Broadcasting is a method that NumPy uses to **perform arithmetic operations on arrays of different shapes**. Instead of manually repeating values to match dimensions, NumPy automatically "broadcasts" smaller arrays across larger arrays so that they are compatible for element-wise operations.

Think of it as NumPy **stretching the smaller array** without making actual copies in memory → which is **fast and memory-efficient**.

### Rules of Broadcasting

For two arrays to be broadcastable:

1. Compare their shapes starting from the **trailing dimensions**.
2. Two dimensions are compatible if:
  - They are equal, OR
  - One of them is 1.

If all dimensions match under these rules, broadcasting is possible.

## ✓ Examples of Broadcasting Rules:

- (4, 3) and (3,) → broadcast → (4, 3)
- (5, 1) and (1, 6) → broadcast → (5, 6)
- (2, 3, 4) and (3, 1) → broadcast → (2, 3, 4)
- ✗ (2, 3) and (4,) → Not compatible

## Examples

### 1. Broadcasting with Scalars

```
import numpy as np

arr = np.array([1, 2, 3, 4, 5])
print(arr + 10)  # Add scalar to array
print(arr * 2)   # Multiply array by scalar
```

Output:

```
[11 12 13 14 15]
[ 2  4  6  8 10]
```

### 2. Broadcasting with Vectors

```
a = np.array([1, 2, 3])      # Shape (3,)
b = np.array([[10], [20], [30]]) # Shape (3,1)

print(a + b)
```

Output:

```
[[11 12 13]
 [21 22 23]
 [31 32 33]]
```

👉 NumPy stretched both arrays to shape (3,3) .

### 3. Broadcasting with Matrices

```
matrix = np.array([[1, 2, 3],
                   [4, 5, 6]]) # Shape (2, 3)

vector = np.array([10, 20, 30]) # Shape (3,)

print(matrix + vector)
```

Output:

```
[[11 22 33]
 [14 25 36]]
```

👉 vector is broadcast across each row of the matrix .

### 4. Complex Example

```
A = np.ones((3, 4)) # Shape (3, 4)
B = np.arange(4)     # Shape (4,)

print(A + B)
```

**Output:**

```
[[1. 2. 3. 4.]
 [1. 2. 3. 4.]
 [1. 2. 3. 4.]]
```

**Tips & Tricks**

- Broadcasting avoids explicit loops → **much faster** than Python loops.
- Use `np.newaxis` (or `None`) to reshape arrays for broadcasting:

```
a = np.array([1, 2, 3]) # (3,)
b = np.array([4, 5, 6]) # (3,)

print(a[:, np.newaxis] + b) # Shape (3,1) + (3,) -> (3,3)
```

**Exercises**

1. Create a 1D array `[1,2,3,4]` and broadcast it to add `100` to each element.
2. Add a column vector `[1,2,3]` to a row vector `[10,20,30]` using broadcasting.
3. Given `A = np.array([[1],[2],[3]])` and `B = np.array([10,20,30])`, apply broadcasting to compute `A * B`.
4. Use broadcasting to normalize a matrix `M` (subtract mean of each column).

**Solutions**

```
# 1
arr = np.array([1,2,3,4])
print(arr + 100)
# [101 102 103 104]

# 2
col = np.array([[1],[2],[3]])
row = np.array([10,20,30])
print(col + row)
# [[11 21 31]
#  [12 22 32]
#  [13 23 33]]

# 3
A = np.array([[1],[2],[3]])
B = np.array([10,20,30])
print(A * B)
# [[10 20 30]
#  [20 40 60]
#  [30 60 90]]

# 4 Normalization by column
M = np.array([[1,2,3],
              [4,5,6],
              [7,8,9]])
col_mean = M.mean(axis=0) # shape (3,)
M_normalized = M - col_mean
print(M_normalized)
# [[-3. -3. -3.]
#  [ 0.  0.  0.]
#  [ 3.  3.  3.]]
```

Start coding or [generate](#) with AI.

## ✓ 7. Linear Algebra with NumPy

NumPy provides a powerful `numpy.linalg` module for linear algebra operations. These are widely used in data science, machine learning, physics, engineering, and mathematics.

### 7.1 Dot Product

The **dot product** of two vectors is the sum of the products of their corresponding elements.

◆ Formula:

$$a \cdot b = \sum_{i=1}^n a_i b_i$$

Example:

```
import numpy as np

a = np.array([1, 2, 3])
b = np.array([4, 5, 6])

dot1 = np.dot(a, b)    # Method 1
dot2 = a @ b           # Method 2 (Python 3.5+)

print(dot1) # 32
print(dot2) # 32
```

✓ Tip: Use @ operator for cleaner code.

### 7.2 Matrix Multiplication

Matrix multiplication is different from element-wise multiplication.

◆ Formula: If  $A$  is of shape  $(m, n)$  and  $B$  is  $(n, p)$ , then result is  $(m, p)$ .

Example:

```
A = np.array([[1, 2], [3, 4]])
B = np.array([[5, 6], [7, 8]])

C1 = np.matmul(A, B)    # Method 1
C2 = A @ B              # Method 2

print(C1)
# [[19 22]
#  [43 50]]
```

### 7.3 Transpose of a Matrix

Transpose flips a matrix over its diagonal (rows become columns).

```
A = np.array([[1, 2, 3], [4, 5, 6]])
print("Original:\n", A)
print("Transpose:\n", A.T)
```

### 7.4 Determinant

The determinant is a scalar value that provides information about a matrix (invertibility, volume scaling).

```
A = np.array([[1, 2], [3, 4]])
det = np.linalg.det(A)
```

```
print("Determinant:", det) # -2.0
```

✓ Tip: If determinant = 0 → Matrix is singular (non-invertible).

## 7.5 Inverse of a Matrix

The inverse of a matrix  $A^{-1}$  satisfies:

$$A \cdot A^{-1} = I$$

```
A = np.array([[1, 2], [3, 4]])
inv = np.linalg.inv(A)
print("Inverse:\n", inv)

# Verification
print("A @ inv:\n", A @ inv) # Should be Identity Matrix
```

## 7.6 Eigenvalues & Eigenvectors

For a square matrix  $A$ :

$$Av = \lambda v$$

Where  $\lambda$  = eigenvalue,  $v$  = eigenvector.

```
A = np.array([[4, -2],
              [1, 1]])

eigenvalues, eigenvectors = np.linalg.eig(A)

print("Eigenvalues:", eigenvalues)
print("Eigenvectors:\n", eigenvectors)
```

## Exercises

Q1. Compute the dot product of  $[2, 3, 4]$  and  $[1, 0, -1]$ .

Q2. Multiply

$$A = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}, B = \begin{bmatrix} 2 & 0 \\ 1 & 2 \end{bmatrix}$$

Q3. Find the determinant of

$$\begin{bmatrix} 6 & 1 \\ 4 & -2 \end{bmatrix}$$

Q4. Find the inverse of

$$\begin{bmatrix} 2 & 1 \\ 7 & 4 \end{bmatrix}$$

Q5. Compute the eigenvalues and eigenvectors of

$$\begin{bmatrix} 3 & 1 \\ 0 & 2 \end{bmatrix}$$

## ✓ Solutions

```
# Q1
a = np.array([2, 3, 4])
b = np.array([1, 0, -1])
print(np.dot(a, b)) # -2
```



```
# Q2
A = np.array([[1, 2], [3, 4]])
B = np.array([[2, 0], [1, 2]])
print(A @ B)
# [[ 4  4]
#  [10  8]]

# Q3
A = np.array([[6, 1], [4, -2]])
print(np.linalg.det(A)) # -16.0

# Q4
A = np.array([[2, 1], [7, 4]])
print(np.linalg.inv(A))
# [[ 4. -1.]
#  [-7.  2.]]

# Q5
A = np.array([[3, 1], [0, 2]])
vals, vecs = np.linalg.eig(A)
print("Eigenvalues:", vals)      # [3.  2.]
print("Eigenvectors:\n", vecs)  # [[1.  0.]
#                               # [0.  1.]]
```

#### ⚡ Tips & Tricks:

- Always check if a matrix is **square & non-singular** before computing its inverse.
- Use `np.allclose(A @ inv, np.eye(n))` to verify inversion accuracy.
- For performance, prefer `np.dot` / `@` over looping through elements.

Start coding or [generate](#) with AI.

## ✓ 8. Random Module in NumPy

NumPy provides powerful tools for generating random numbers, which are essential in simulations, testing, and machine learning. These random functions are available in `numpy.random`.

### 8.1 Generating Random Numbers

#### 1. `np.random.rand()`

- Generates random numbers from a **uniform distribution** between 0 and 1.
- You can specify shape as arguments.

```
import numpy as np

# Single random number
print(np.random.rand())

# 1D array of 5 random numbers
print(np.random.rand(5))

# 2D array (3x3) of random numbers
print(np.random.rand(3, 3))
```

#### 2. `np.random.randn()`

- Generates random numbers from a **standard normal distribution** (mean=0, variance=1).

```
# Single random number
print(np.random.randn())
```

```
# 1D array of 5 random numbers
print(np.random.randn(5))

# 2D array (2x3) from standard normal distribution
print(np.random.randn(2, 3))
```

### 3. np.random.randint()

- Generates random **integers** in a given range `[low, high)`.

```
# Single integer between 1 and 10
print(np.random.randint(1, 10))

# Array of 5 integers between 0 and 50
print(np.random.randint(0, 50, 5))

# 2D array of random integers between 10 and 100
print(np.random.randint(10, 100, size=(3, 4)))
```

## 8.2 Random Choice

Select random elements from a sequence.

```
arr = np.array([10, 20, 30, 40, 50])

# Pick one random element
print(np.random.choice(arr))

# Pick 3 random elements (with replacement)
print(np.random.choice(arr, 3))

# Pick 3 random elements (without replacement)
print(np.random.choice(arr, 3, replace=False))

# Probability distribution
print(np.random.choice(arr, 5, p=[0.1, 0.2, 0.3, 0.1, 0.3]))
```

## 8.3 Seeding Random Numbers

Random numbers are generated based on a seed value. If you set a seed, the random numbers will always be the **same**, which is useful for reproducibility in experiments.

```
np.random.seed(42)

print(np.random.rand(3))
print(np.random.randint(1, 10, 3))

# Resetting seed gives the same result
np.random.seed(42)
print(np.random.rand(3))
```

### Exercises

- Generate a **3x3 matrix** of random integers between **1 and 100**.
- Create an array of **10 random numbers** from a **normal distribution**.
- Use `np.random.choice` to pick **5 random students** from the list:

```
students = ["Ali", "Sara", "John", "Mary", "Ahmed", "Zara", "Omar", "Fatima"]
```

4. Set a seed and generate **two identical arrays** of random numbers.

## Solutions

```
# 1. 3x3 random integers
mat = np.random.randint(1, 100, size=(3, 3))
print("3x3 Random Integers:\n", mat)

# 2. 10 numbers from normal distribution
normal_nums = np.random.randn(10)
print("Normal Distribution:\n", normal_nums)

# 3. Random choice from list
students = ["Ali", "Sara", "John", "Mary", "Ahmed", "Zara", "Omar", "Fatima"]
picked = np.random.choice(students, 5, replace=False)
print("Picked students:", picked)

# 4. Seeding
np.random.seed(7)
arr1 = np.random.rand(5)

np.random.seed(7)
arr2 = np.random.rand(5)

print("Array 1:", arr1)
print("Array 2:", arr2)
```

## Tips & Tricks

- Use `np.random.seed()` to ensure **reproducibility**.
- Use `replace=False` in `np.random.choice()` for **sampling without repetition**.
- For **machine learning experiments**, always fix the seed before data splitting or weight initialization.

Start coding or [generate](#) with AI.

## ✓ 9. Working with Files in NumPy

Often, after processing or generating data, we need to save it for later use. NumPy provides efficient ways to store and retrieve arrays.

### 9.1 Saving Arrays

#### 1. `np.save()`

- Saves an array in **binary format** (`.npy`).
- Best for saving and loading arrays quickly without losing precision.

```
import numpy as np

arr = np.array([10, 20, 30, 40, 50])

# Save array
np.save("my_array.npy", arr)

print("Array saved successfully!")
```

#### 2. `np.savetxt()`

- Saves an array in **text format** (`.txt` or `.csv`).
- Can specify delimiters (`,` or `\t`).
- Useful when you want to **share data in a readable format**.

```
arr2d = np.array([[1, 2, 3], [4, 5, 6]])

# Save as text file
np.savetxt("my_array.txt", arr2d)

# Save as CSV with delimiter
np.savetxt("my_array.csv", arr2d, delimiter=",")
```

## 9.2 Loading Arrays

### 1. np.load()

- Loads arrays stored in `.npy` or `.npz` files.

```
loaded = np.load("my_array.npy")
print("Loaded Array:", loaded)
```

### 2. np.loadtxt()

- Loads arrays stored in **text files** (`.txt`, `.csv`).

```
loaded_txt = np.loadtxt("my_array.txt")
print("Loaded from TXT:", loaded_txt)

loaded_csv = np.loadtxt("my_array.csv", delimiter=",")
print("Loaded from CSV:\n", loaded_csv)
```

### 3. np.genfromtxt()

- Similar to `np.loadtxt()` but **handles missing values**.
- Useful when files have incomplete data.

```
# Example CSV with missing values
# 1,2,3
# 4,,6
# 7,8,9

loaded_gen = np.genfromtxt("my_incomplete.csv", delimiter=",", filling_values=-1)
print("Loaded with genfromtxt:\n", loaded_gen)
```

## Exercises

1. Save a **2D array** into a `.npy` file and load it back.
2. Save an array into a **CSV file** with commas as delimiters, then load it.
3. Load a text file with missing values using **np.genfromtxt** and replace missing values with **0**.

## Solutions

```
# 1. Save and load .npy
arr = np.array([[10, 20, 30], [40, 50, 60]])
np.save("matrix.npy", arr)
loaded = np.load("matrix.npy")
print("Loaded from .npy:\n", loaded)

# 2. Save and load CSV
np.savetxt("matrix.csv", arr, delimiter=",")
loaded_csv = np.loadtxt("matrix.csv", delimiter=",")
print("Loaded from CSV:\n", loaded_csv)

# 3. Handle missing values
```

```
# Assume file 'data_with_missing.csv' contains some empty entries
loaded_gen = np.genfromtxt("data_with_missing.csv", delimiter=",", filling_values=0)
print("Loaded with missing values replaced:\n", loaded_gen)
```

### 🔥 Tips & Tricks

- Use **.numpy files** for fast saving/loading of arrays in projects.
- Use **.csv or .txt** when data needs to be shared with others.
- Use `np.genfromtxt()` when dealing with **messy real-world datasets**.

Start coding or [generate](#) with AI.

## ✓ 10. Advanced Topics in NumPy

### 10.1 Vectorization vs. Loops

- NumPy is optimized for **vectorized operations** (using array math instead of Python loops).
- Vectorization uses **C-optimized code**, making it much faster.

💠 Example: Compute the square of numbers from 1 to 1,000,000

```
import numpy as np
import time

# Using loop (slow)
arr = np.arange(1, 1000001)
start = time.time()
result_loop = [x**2 for x in arr]
end = time.time()
print("Loop time:", end - start)

# Using vectorization (fast)
start = time.time()
result_vec = arr ** 2
end = time.time()
print("Vectorized time:", end - start)
```

✅ Vectorization is **10–100x faster** than Python loops.

### 10.2 Memory Views & Copying

- NumPy arrays can **share memory** (views) or create a **new copy**.
- Views are more memory-efficient.

```
arr = np.array([10, 20, 30, 40])

# Create a view (changes affect original)
view = arr.view()
view[0] = 99
print("Original after view change:", arr)

# Create a copy (independent array)
copy = arr.copy()
copy[1] = 77
print("Original after copy change:", arr)
```

✅ Use `.copy()` when you need an independent object.

### 10.3 Structured Arrays & dtypes

NumPy supports **structured arrays**, where each element can have multiple fields (like database rows).

```
# Define structured dtype
student_dtype = np.dtype([
    ('name', 'U10'),    # string up to 10 chars
    ('age', 'i4'),      # 4-byte integer
    ('marks', 'f4')     # 4-byte float
])

# Create structured array
students = np.array([
    ("Alice", 20, 85.5),
    ("Bob", 22, 90.0),
    ("Charlie", 21, 78.0)
], dtype=student_dtype)

print(students)
print("Names:", students['name'])
print("Marks:", students['marks'])
```

✔ Structured arrays are useful for datasets with **mixed data types**.

## 10.4 Masking & Conditional Selection

- Boolean masking allows **filtering arrays**.
- Similar to SQL queries but inside NumPy.

```
arr = np.array([5, 12, 18, 7, 30, 22])

# Get elements greater than 15
mask = arr > 15
print("Mask:", mask)
print("Filtered:", arr[mask])

# Direct condition
print("Even numbers:", arr[arr % 2 == 0])
```

✔ Very useful for **data cleaning and filtering**.

## 10.5 Stacking & Splitting Arrays

### Stacking

Combine multiple arrays.

```
a = np.array([1, 2, 3])
b = np.array([4, 5, 6])

# Horizontal stacking
print("hstack:", np.hstack((a, b)))

# Vertical stacking
print("vstack:\n", np.vstack((a, b)))
```

### Splitting

Divide arrays into multiple parts.

```
arr = np.array([10, 20, 30, 40, 50, 60])

# Split into 3 equal parts
parts = np.split(arr, 3)
print("Split into 3 parts:", parts)
```

```
# 2D split
matrix = np.array([[1, 2], [3, 4], [5, 6]])
upper, lower = np.vsplit(matrix, 2)
print("Upper:\n", upper)
print("Lower:\n", lower)
```

## ✓ Quick Recap

1. **Vectorization** → Faster than loops.
2. **Copy vs View** → `.copy()` makes new memory, `.view()` shares memory.
3. **Structured Arrays** → Store mixed data types.
4. **Masking** → Filter arrays with conditions.
5. **Stacking & Splitting** → Combine or divide arrays.

Start coding or [generate](#) with AI.

## ✓ NumPy Complete Tutorial (Beginner to Advanced)

This guide covers **all essential and advanced topics in NumPy** with examples and explanations. Use it as a reference or a learning roadmap.

## Table of Contents

1. Introduction to NumPy
2. NumPy Arrays (Creation & Basics)
3. Array Indexing & Slicing
4. Array Operations
5. Universal Functions (ufuncs)
6. Broadcasting
7. Mathematical & Statistical Functions
8. Linear Algebra
9. Random Module in NumPy
10. Advanced Topics

## 1. Introduction to NumPy

- NumPy = Numerical Python → powerful library for numerical computing.
- Provides **N-dimensional arrays** (ndarrays).
- Faster than Python lists because operations are vectorized and written in C.

```
import numpy as np
print(np.__version__)
```

## 2. NumPy Arrays (Creation & Basics)

```
# Create array from list
arr = np.array([1, 2, 3])
print(arr)

# Special arrays
zeros = np.zeros((2, 3))      # 2x3 matrix of zeros
ones = np.ones((3, 3))       # 3x3 matrix of ones
arange = np.arange(0, 10, 2)  # [0,2,4,6,8]
linspace = np.linspace(0, 1, 5) # 5 numbers between 0 and 1
```

- Attributes:

```
print(arr.ndim) # dimensions
print(arr.shape) # shape
print(arr.size) # total elements
print(arr.dtype) # data type
```

---

### 3. Array Indexing & Slicing

```
arr = np.array([10, 20, 30, 40, 50])
print(arr[0]) # first element
print(arr[-1]) # last element

# Slicing
print(arr[1:4]) # [20,30,40]

# 2D indexing
mat = np.array([[1,2,3],[4,5,6],[7,8,9]])
print(mat[0,1]) # element at row 0 col 1 → 2
print(mat[:,0]) # first column
```

---

### 4. Array Operations

```
arr1 = np.array([1, 2, 3])
arr2 = np.array([4, 5, 6])

print(arr1 + arr2) # [5,7,9]
print(arr1 * arr2) # [4,10,18]
print(arr1 ** 2) # [1,4,9]
```

---

### 5. Universal Functions (ufuncs)

- Fast element-wise functions.

```
arr = np.array([1, 4, 9, 16])
print(np.sqrt(arr))
print(np.exp(arr))
print(np.log(arr))
```

---

### 6. Broadcasting

- NumPy automatically expands smaller arrays.

```
mat = np.array([[1,2,3],[4,5,6]])
vec = np.array([10,20,30])
print(mat + vec)
```

---

### 7. Mathematical & Statistical Functions

```
arr = np.array([1,2,3,4,5])
print(np.sum(arr))
print(np.mean(arr))
print(np.std(arr))
print(np.min(arr))
print(np.max(arr))
```



## 8. Linear Algebra

```
A = np.array([[1,2],[3,4]])
B = np.array([[5,6],[7,8]])

# Matrix multiplication
print(np.dot(A, B))

# Transpose
print(A.T)

# Inverse
print(np.linalg.inv(A))

# Eigenvalues & vectors
vals, vecs = np.linalg.eig(A)
print(vals)
print(vecs)
```

## 9. Random Module in NumPy

```
np.random.seed(42) # reproducibility

print(np.random.rand(3))    # uniform [0,1)
print(np.random.randn(3))   # normal distribution
print(np.random.randint(1,10,5)) # random integers
```

## 10. Advanced Topics

### 10.1 Vectorization vs Loops

```
import time
arr = np.arange(1, 1000001)

start = time.time()
res = [x**2 for x in arr]
print("Loop time:", time.time()-start)

start = time.time()
res = arr**2
print("Vectorized time:", time.time()-start)
```

### 10.2 Memory Views vs Copies

```
arr = np.array([10,20,30])
view = arr.view()
view[0] = 99
print(arr)

copy = arr.copy()
copy[1] = 77
print(arr)
```

### 10.3 Structured Arrays

```
dt = np.dtype([('name', 'U10'), ('age', 'i4'), ('marks', 'f4')])
students = np.array([
    ("Alice", 20, 85.5),
    ("Bob", 22, 90.0)
```

```
], dtype=dt)

print(students['name'])
print(students['marks'])
```

## 10.4 Masking & Conditional Selection

```
arr = np.array([5,12,18,7,30])
print(arr[arr > 15])
print(arr[arr % 2 == 0])
```

## 10.5 Stacking & Splitting


```
a = np.array([1,2,3])
b = np.array([4,5,6])

print(np.hstack((a,b)))
print(np.vstack((a,b)))

arr = np.array([10,20,30,40,50,60])
print(np.split(arr,3))
```

## Recap

- NumPy is essential for scientific computing.
- Learn **creation, indexing, operations, broadcasting, linear algebra, random module, and advanced memory tricks**.
- Prefer **vectorization over loops** for speed.
- Use **masking & stacking** for clean data manipulation.

 You now have a **complete NumPy roadmap** from basics → advanced. Would you like me to also add a **set of practice exercises with solutions** at the end of this guide (like a mini workbook)?

Start coding or generate with AI.